

# Recurrent Neural Networks (RNNs)

Michel RIVEILL  
[michel.riveill@univ-cotedazur.fr](mailto:michel.riveill@univ-cotedazur.fr)

# Motivation


- ▶ Humans don't start their thinking from scratch every second
  - ▶ Thoughts have persistence
- ▶ Traditional neural networks can't characterize this phenomena
  - ▶ Ex: classify what is happening at every point in a movie
  - ▶ How a neural network can inform later events about the previous ones
- ▶ **Recurrent neural networks** address this issue
  - ▶ Some applications
    - ▶ NER - Naming Entity Recognition
      - Same word may have a different label depending on the context.
        - **Apple** CEO Tim Cook eat an **apple**
    - ▶ Forecasting - Time-series Prediction
- ▶ **How?**
  - ▶ **Add state to artificial neurons**

# What are RNNs?

- ▶ Main idea is to make use of sequential information
- ▶ How RNN is different from neural network?
  - ▶ Vanilla neural networks (MLP) **assume** all inputs and outputs are independent of each other
  - ▶ But for many tasks, that's a very bad idea
- ▶ What RNN does?
  - ▶ Perform the same task for every element of a sequence (that's what **recurrent** stands for)
  - ▶ Output depends on the previous computations!
- ▶ Another way of interpretation – RNNs have a “**memory**”
  - ▶ To store previous computations


# From vanilla NN to recurrent NN

- ▶ Vanilla cell
  - ▶  $y = F(U.X)$




# From vanilla NN to recurrent NN

- ▶ Vanilla cell
  - ▶  $y = F(U.X)$
- ▶ Recurrent cell → use 2 weights matrix
  - ▶ Add an internal variable:  $h$
  - ▶ The output depends to the current entry and the previous internal variable:
    - ▶  $h_t = F(W.h_{t-1} + U.X_t)$
    - ▶ Could be rewritten on  $h_t = F(W.[h_{t-1}, X_t])$



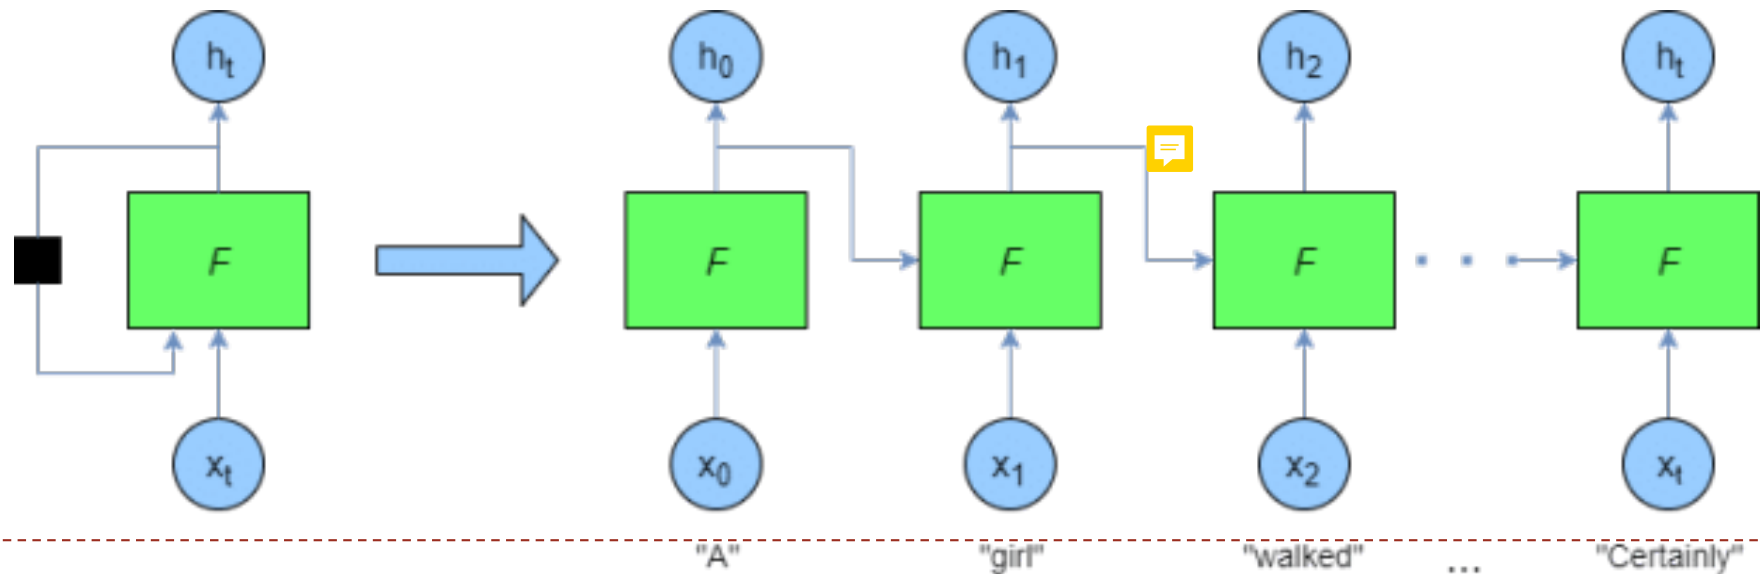
# From vanilla NN to recurrent NN

- ▶ Vanilla cell
  - ▶  $y = F(U.X)$
- ▶ Recurrent cell → use 2 weights matrix
  - ▶  $h_t = F(W.[h_{t-1}, X_t])$
- ▶ Recurrent layer, step by step
  - ▶ at each time step
    - ▶ A new entry is being supplied
    - ▶ And a new output ( $h_t$ ) is calculated using:
      - The new input  $X_t$
      - The output of the previous step  $h_{t-1}$
  - ▶  $h_1 = F(W.[h_0, X_1])$
  - ▶  $h_2 = F(W.[h_1, X_2])$
  - ▶  $h_3 = F(W.[h_2, X_3])$
  - ▶ ...




# From vanilla NN to recurrent NN

- ▶ Vanilla cell
  - ▶  $y = F(U.X)$
- ▶ Recurrent cell
  - ▶  $h_t = F(W.[h_{t-1}, X_t])$
- ▶ Recurrent neural networks are “unrolled” programmatically during training and prediction
  - ▶ All neurons share the same weight matrix




# Remember

From




To






# Remember


**From**



**To**

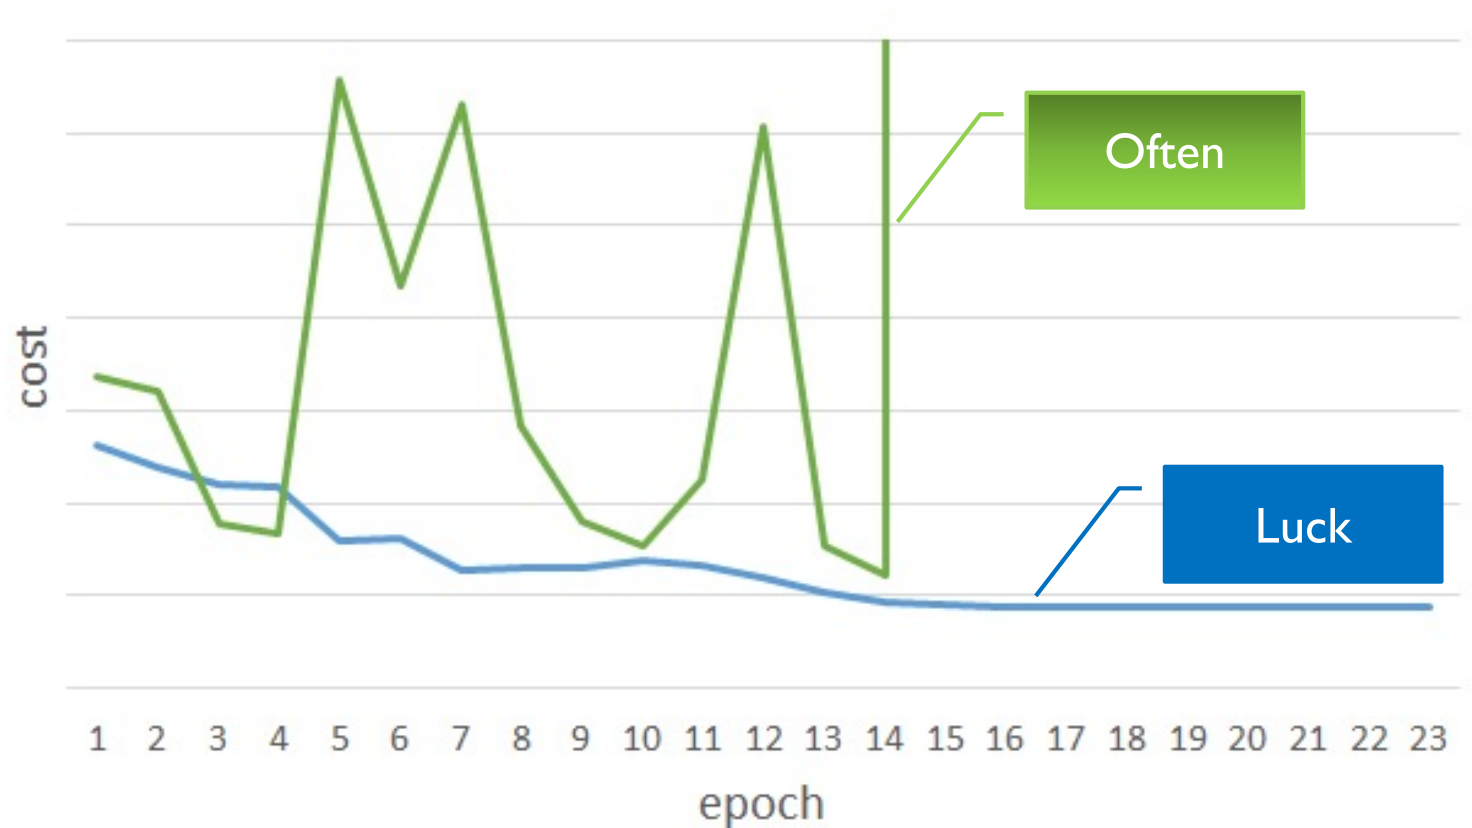


# RNN in action



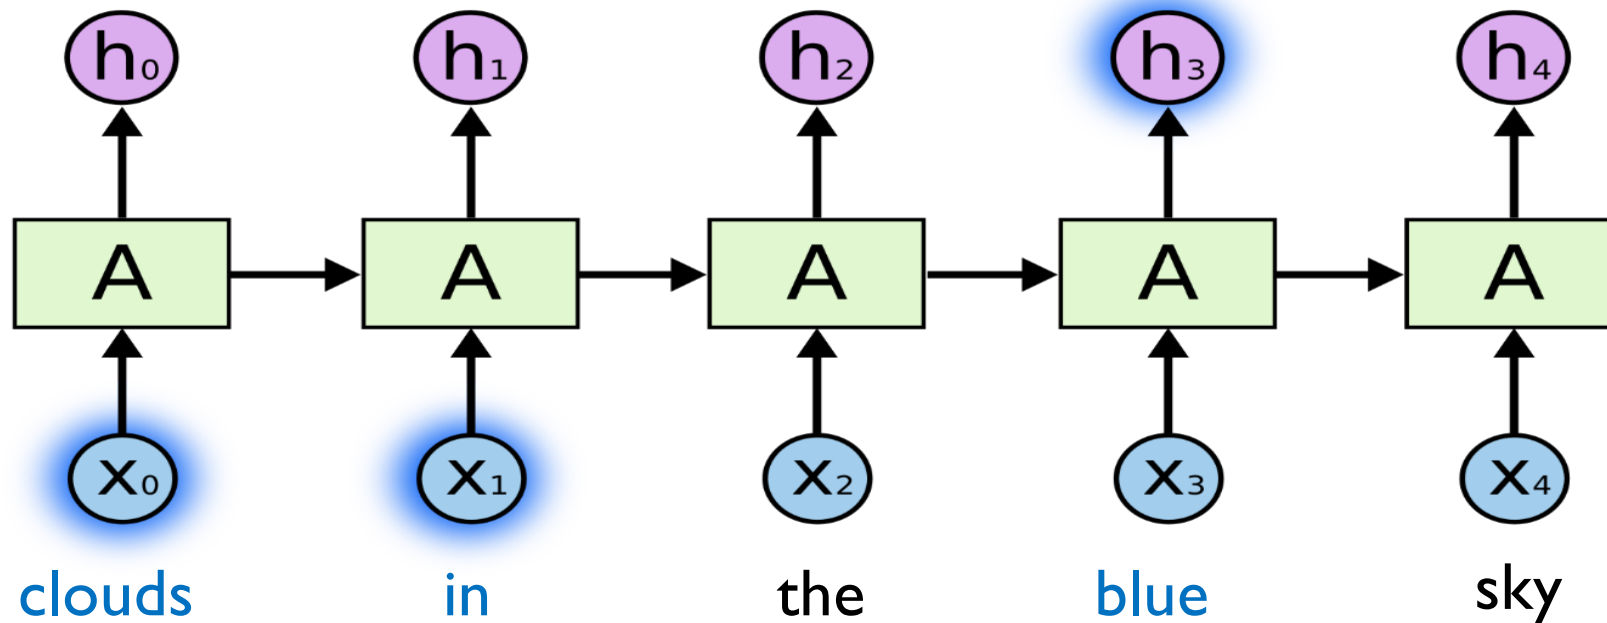
# Problems with naive RNN

- ▶ RNNs do not learn easily
- ▶ Unfolding the network for learning leads to vanishing gradient problems!




# From vanilla RNN...

- ▶ The context is close to the word to be predicted
  - ▶ Few iterations separate them.
  - ▶ No problem



## ... to LSTM (Long Short Term Memory)


- ▶ The context is far from the word to predict
  - ▶ Many iterations separate them!
  - ▶ Possible gradient problem




I grew up in France...


I speak French...

# From Vanilla to LSTM Cells

- ▶ It is necessary to prevent the gradient from disappearing...
- ▶ Normally, the network memory is 
  - ▶  $h_t = \tanh(W \square [h_{t-1}, x_t])$
  - ▶ Involves a single level of processing
  - ▶ Creating the risk of the evanescent gradient.



# Dealing with the vanishing gradient problem → LSTM cell




*LSTM = Long Short-Term Memory*

*(crédit : image modifiée de Michaël Nguyen)*

# LSTM cell


- ▶ Cellule composée de trois “portes” : ce sont des zones de calculs qui régulent le flot d’informations (en réalisant des actions spécifiques)
  - ▶ Forget gate (porte d’oubli)
  - ▶ Input gate (porte d’entrée)
  - ▶ Output gate (porte de sortie)
- ▶ Hidden state (état caché)
- ▶ Cell state (état de la cellule)
  - ▶ Like residual




(crédit : image modifiée de Michaël Nguyen)



# LSTM cell (porte oubli / forget get)




# LSTM cell (porte entrée / input get)




Capture l'information d'entrée qui doit être incluse dans l'état de la cellule

# LSTM cell (état de la cellule / cell state)



L'état de la cellule se calcule assez simplement à partir de la porte d'oubli et de la porte d'entrée

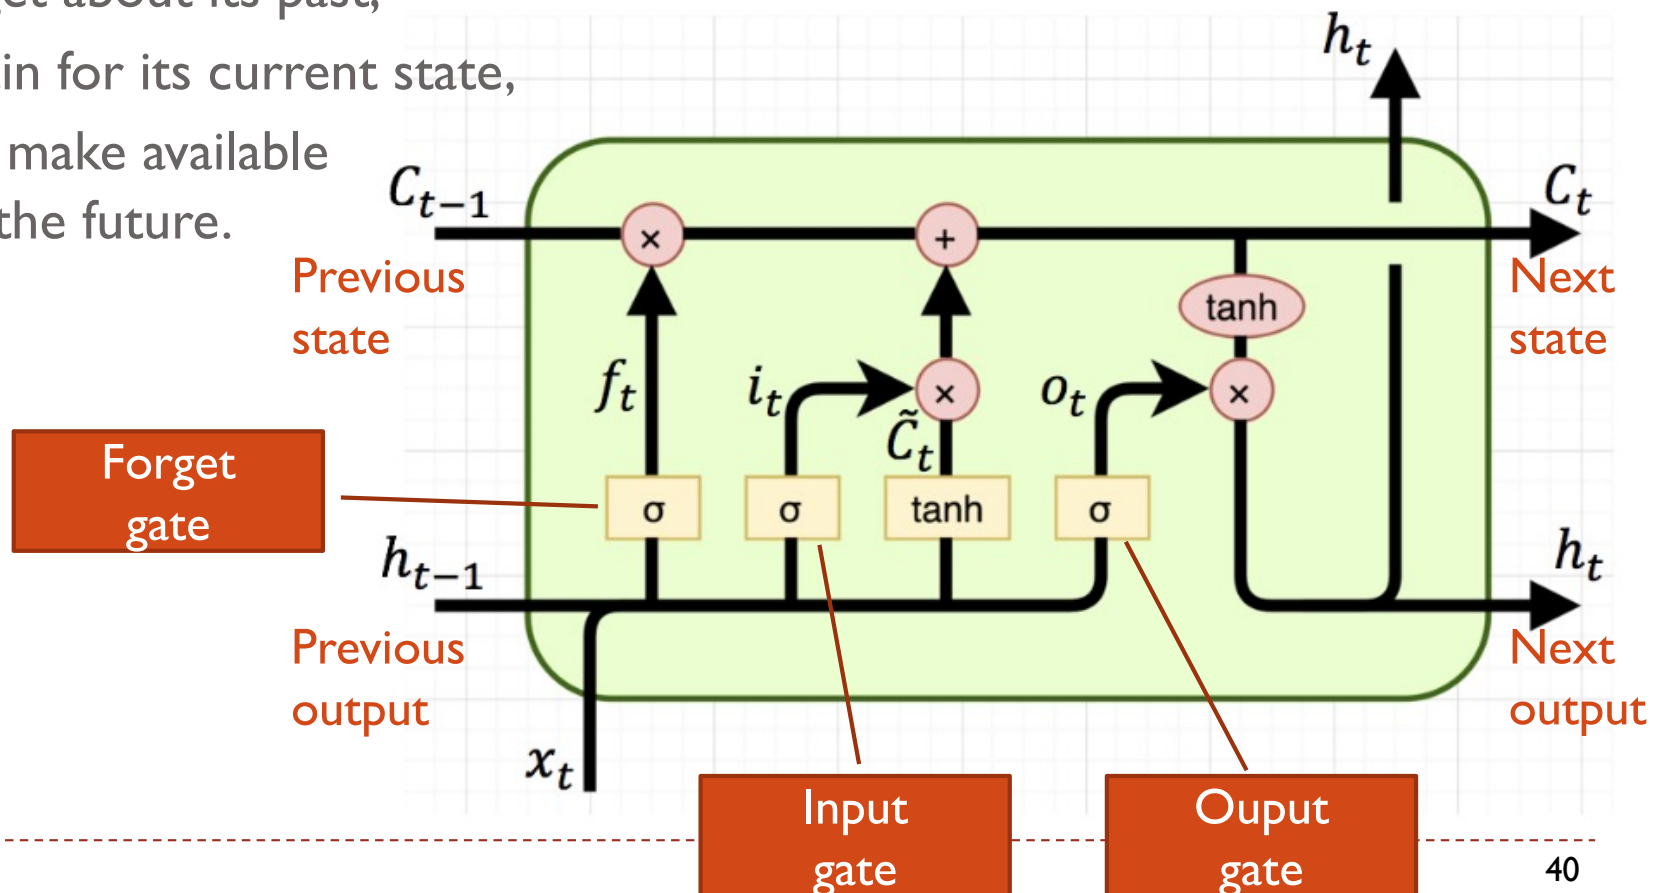
# LSTM cell (porte de sortie / output gate)



La porte de sortie décide quel sera le prochain état caché. Il contient des informations sur les entrées précédentes du réseau et sert aux prédictions.

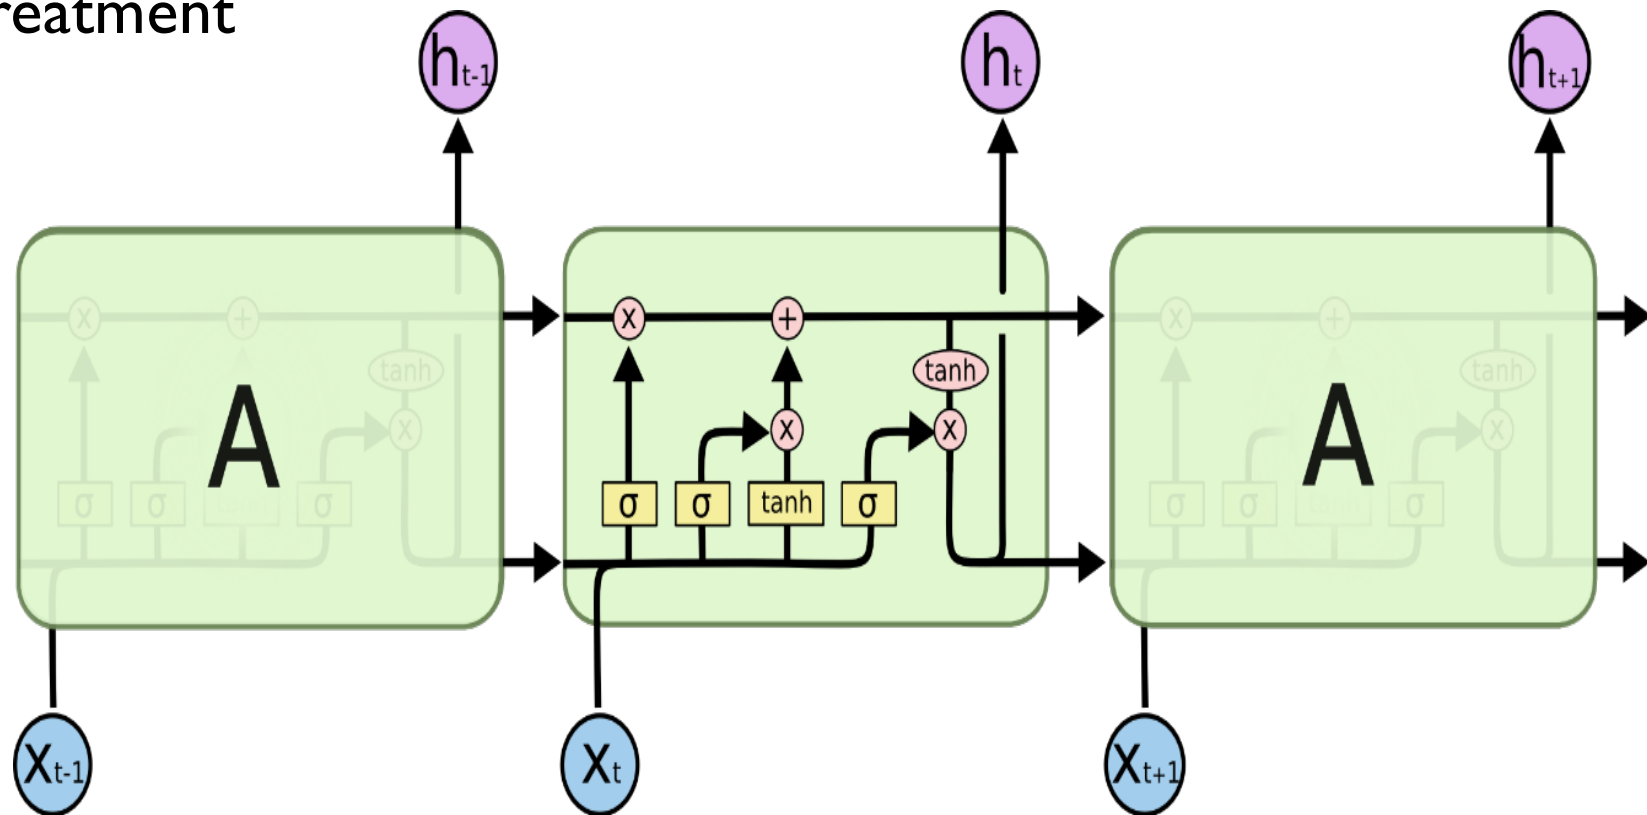
# LSTM Cells

- ▶ Adds a context memory that affects the information flow and its processing (cell state).
- ▶ Three gates decide what a cell should
  - ▶ forget about its past,
  - ▶ retain for its current state,
  - ▶ and make available for the future.



# LSTM Cells

- Concretely, recurrence in an LSTM cell involves 4 levels of treatment



Neural Network  
Layer

Pointwise  
Operation

Vector  
Transfer

Concatenate

Copy



# Naïve RNN vs LSTM

## ▶ Naïve RNN

- ▶ Reuse at each step the previous output

## ▶ LSTM

- ▶ At each step 3 gate control the use of Input value, Cell state and previous output




c changes slowly →  $c^t$  is  $c^{t-1}$  added by something

h changes faster →  $h^t$  and  $h^{t-1}$  can be very different

# GRU – gated recurrent unit

- ▶ GRU = a light LSTM Cell



$$z_t = \sigma (W_z \cdot [h_{t-1}, x_t])$$

$$r_t = \sigma (W_r \cdot [h_{t-1}, x_t])$$

$$\tilde{h}_t = \tanh (W \cdot [r_t * h_{t-1}, x_t])$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t$$

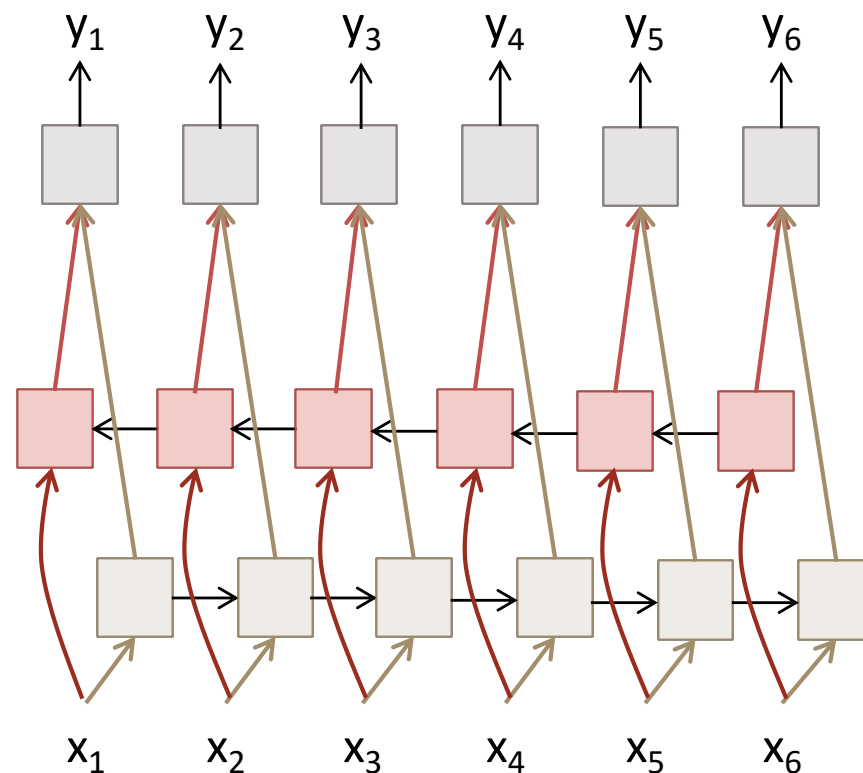
- It combines the **forget** and **input** into a single **update gate**.
  - It also **merges** the **cell state** and **hidden state**.
- This is simpler than LSTM.




# Bi-directional RNNs

---

- ▶ RNNs can process the input sequence in forward and in the reverse direction




- Popular in speech recognition, could be used also with text




## RNN cell in Keras




# Keras Long Short-Term Memory Cell

- ▶ **from tf.keras.layers import LSTM**
- ▶ Main params
  - ▶ **Units:** dimension of output space
  - ▶ **return\_sequences:** True or False 
    - ▶ If False return only the last output
    - ▶ If True return the full sequence of the output sequence
      - Output sequence = hidden state (the vocabulary change regarding documentation)
  - ▶ **return\_state:** True or False
    - ▶ If True return 3 values
      - The full output sequence or only the last one (depend on return\_sequences)
      - The last output sequence
      - The cell state
    - ▶ If False return nothing
  - ▶ **stateful:** True or False
    - ▶ If True, the last state for each sample at index i in a batch will be used as initial state for the sample of index i in the following batch.
    - ▶ You have to put **shuffle=False** in a fit method




*return\_sequences = False*



*return\_sequences = True*

with `return_sequences=False`,  
**Dense** layer is applied only  
once at the last cell

If `return_sequences=True`  
**Dense** layer is applied to every  
timestep



# Keras Gated Recurrent Unit Cell

- ▶ **from keras.layers import GRU**
- ▶ Main params (**similar to LSTM**)
  - ▶ **Units:** dimension of output space
  - ▶ **return\_sequences:** True or False
    - ▶ If False return only the last output
    - ▶ If True return the full sequence of the output sequence
      - Output sequence = hidden state (the vocabulary change regarding documentation)
  - ▶ **return\_state:** True or False
    - ▶ If True return 3 values
      - The full output sequence or only the last one (depend on return\_sequences)
      - The last output sequence
      - The cell state
    - ▶ If False return nothing
  - ▶ **stateful:** True or False
    - ▶ If True, the last state for each sample at index i in a batch will be used as initial state for the sample of index i in the following batch.
    - ▶ **You have to put shuffle=False in a fit method**

# A basic example

- ▶ `inputs = Input(shape=(SEQUENCE_SIZE,))`
- ▶ `embedding = Embedding(VOCABULARY_SIZE,  
 EMBEDDING_SIZE,  
 input_length=SEQUENCE_SIZE)(inputs)`
- ▶ `output = LSTM(16, return_sequences=False,  
 activation='relu')(embedding)`
- ▶ `predictions = Dense(nb_classes,  
 activation='softmax')(output)`
- ▶ **Fit by batch**
  - ▶ `Model.fit(X, y, ...)`. ← all item have the same length
- ▶ **Fit by item**
  - ▶ For `i in range(len(X))`: ← could be different length
    - ▶ `Model.fit(X[i], y[i], ...)`

# Some use of RNN

## → Text Classification / Sentiment analysis

### Affect a label to a text

- ▶ Classify a

- ▶ restaurant review from Yelp!

- ▶ movie review from IMDB

- ...

- as positive or negative

- ▶ Inputs:

- ▶ Multiple words, one or more sentences

- ▶ Outputs:


- ▶ Positive / Negative classification

- ▶ “The food was really good”


- ▶ “The chicken crossed the road because it was uncooked”




# Sentiment analysis



# Sentiment analysis

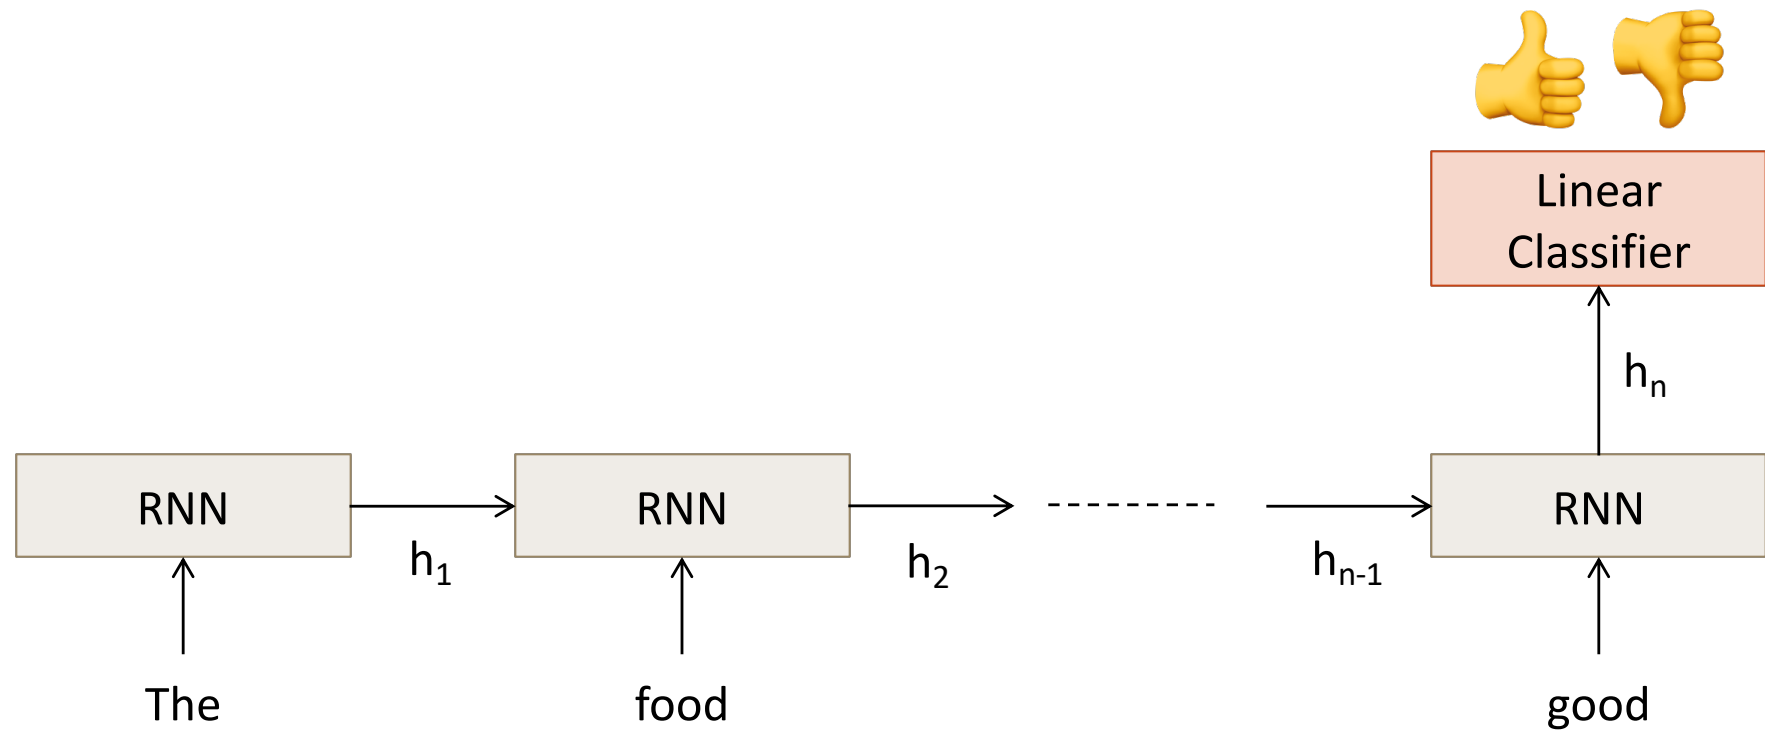


# Sentiment analysis



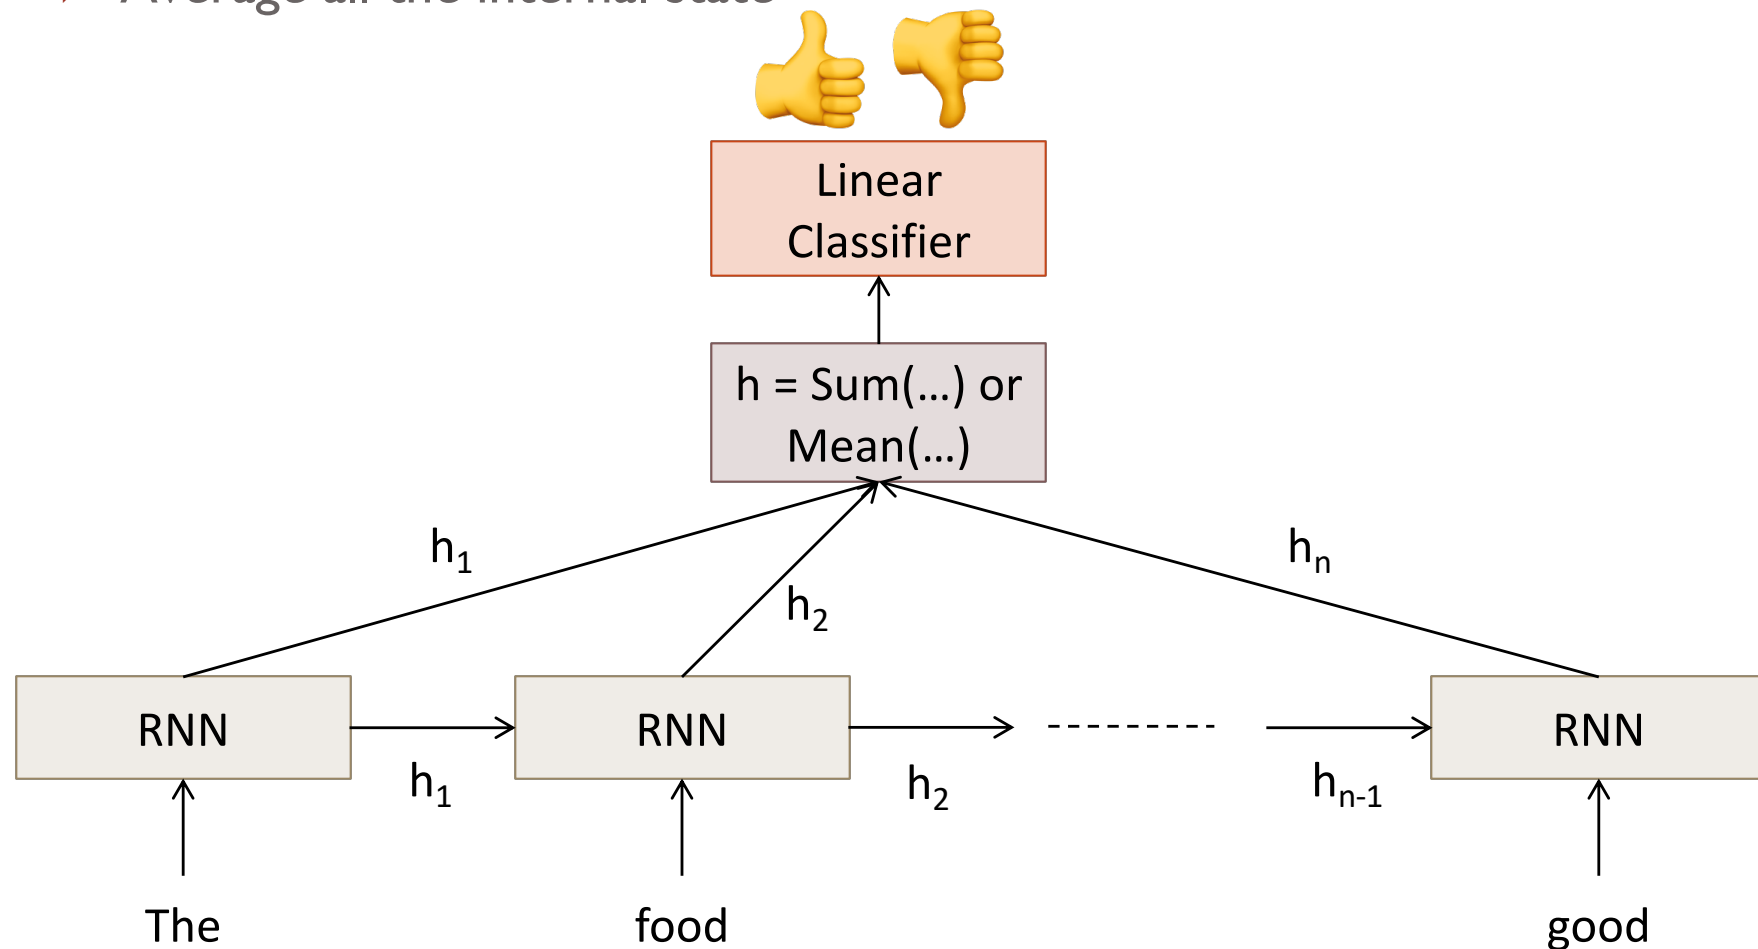
# Sentiment analysis - solution 1

- ▶ retrieve only the last state



# Sentiment analysis – solution2

- ▶ Other possible architecture
  - ▶ Average all the internal state



# Some use of RNN

## → Named Entity Recognition / Part of Speech Tagging

- ▶ Affect a label to each word
  - ▶ **find** and **classify** names in text
    - ▶ Could be an entity : number, country, person, ... (NER)
    - ▶ Could be a function : noun, verb, adverbs, ... (POS)

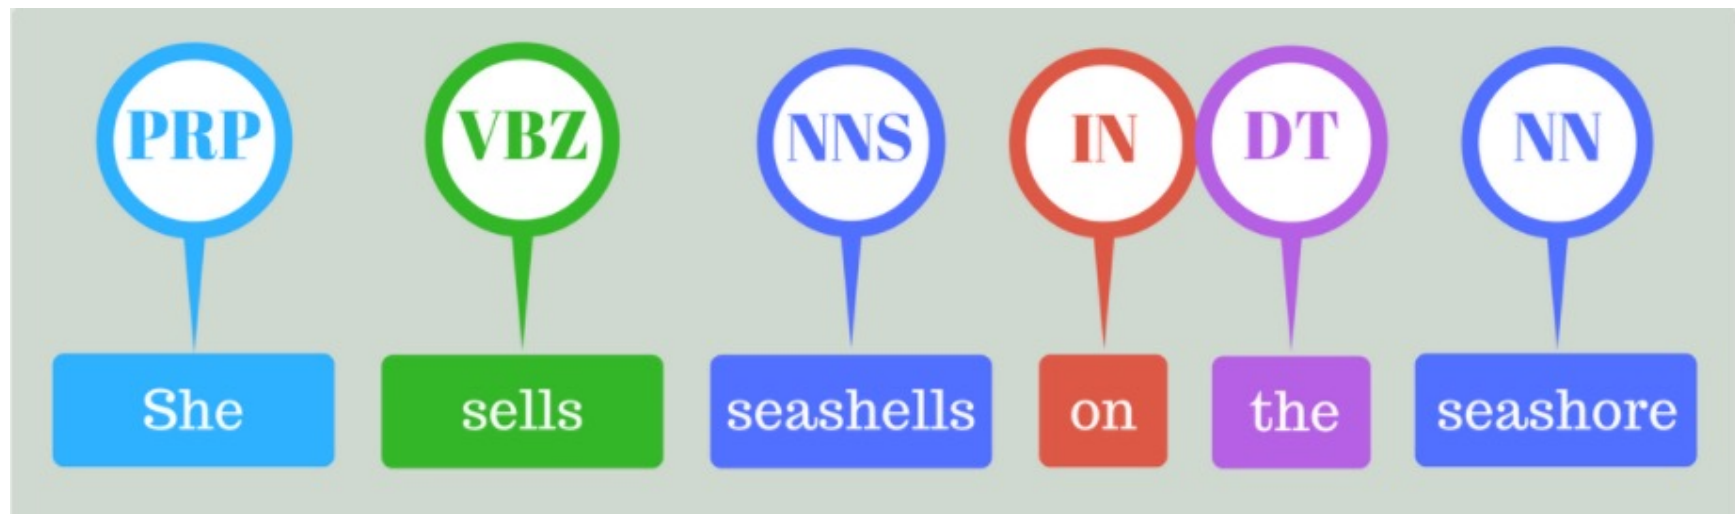
In fact, the **Chinese** **NORP** market has the **three** **CARDINAL** most influential names of the retail and tech space – **Alibaba** **GPE**, **Baidu** **ORG**, and **Tencent** **PERSON** (collectively touted as **BAT** **ORG**), and is betting big in the global **AI** **GPE** in retail industry space. The **three** **CARDINAL** giants which are claimed to have a cut-throat competition with the **U.S.** **GPE** (in terms of resources and capital) are positioning themselves to become the 'future **AI** **PERSON** platforms'. The trio is also expanding in other **Asian** **NORP** countries and investing heavily in the **U.S.** **GPE** based **AI** **GPE** startups to leverage the power of **AI** **GPE**. Backed by such powerful initiatives and presence of these conglomerates, the market in APAC AI is forecast to be the fastest-growing **one** **CARDINAL**, with an anticipated **CAGR** **PERSON** of **45%** **PERCENT** over **2018 - 2024** **DATE**.

To further elaborate on the geographical trends, **North America** **LOC** has procured **more than 50%** **PERCENT** of the global share in **2017** **DATE** and has been leading the regional landscape of **AI** **GPE** in the retail market. The **U.S.** **GPE** has a significant credit in the regional trends with **over 65%** **PERCENT** of investments (including M&As, private equity, and venture capital) in artificial intelligence technology. Additionally, the region is a huge hub for startups in tandem with the presence of tech titans, such as **Google** **ORG**, **IBM** **ORG**, and **Microsoft** **ORG**.

# Some use of RNN

## → Named Entity Recognition / Part of Speech Tagging

- ▶ Affect a label to each word
  - ▶ **find** and **classify** names in text
    - ▶ Could be an entity : number, country, person, ... (NER)
    - ▶ Could be a function : noun, verb, adverbs, ... (POS)




# Label representation – BIO tags

- ▶ Labels can be for words or groups of words
  - ▶ Adam Smith works for IBM , London .
- ▶ To represent this, a "BIO" representation is generally used.
  - ▶ B beginning of an entity
  - ▶ I continues the entity
  - ▶ O word outside the entity
- ▶ For example
  - ▶ ['Adam', 'Smith', 'works', 'for', 'IBM', ',', 'London', '.']
  - ▶ Without BIO: [PER, PER, O, O, ORG, O, GEO, O]
  - ▶ With BIO: [B\_PER, I\_PER, O, O, B\_ORG, O, B\_GEO, O]




# MLP for NER




- Activation function for output: softmax
- Labels are OneHotEncoded

# Recurrent neural network for NER

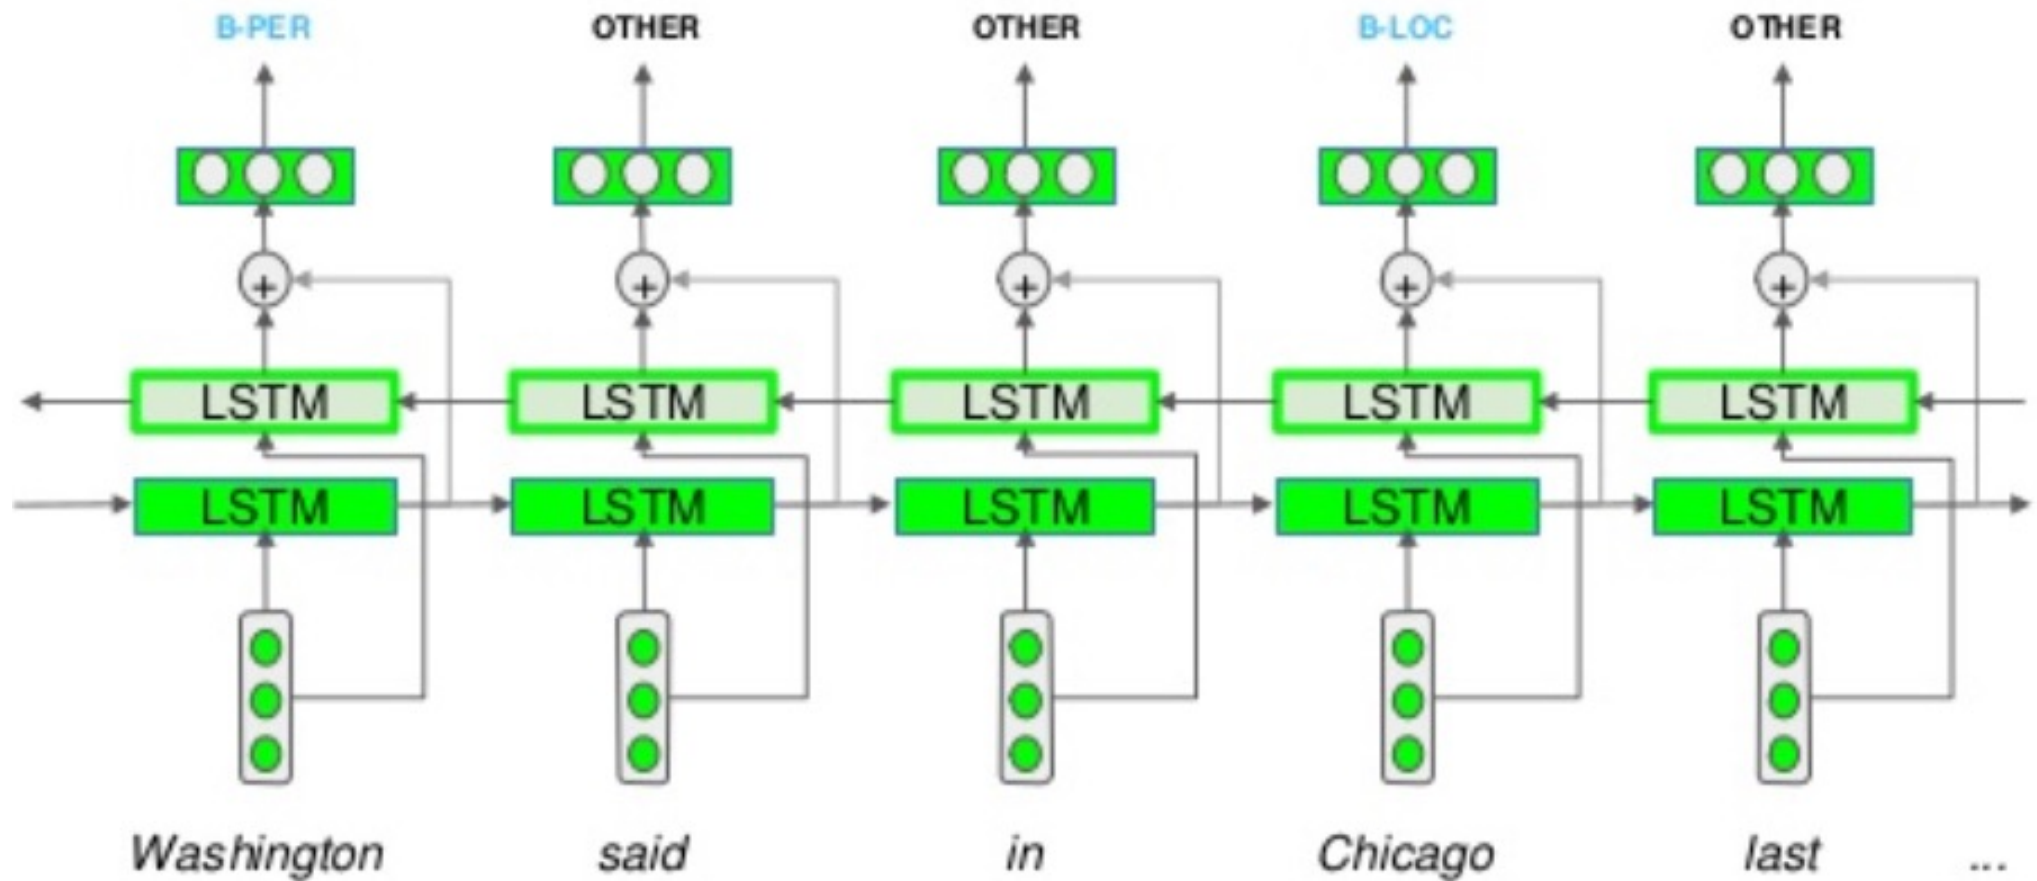


# Recurrent neural network for NER (unfolded)




- Activation function for output: softmax
- Labels are OneHotEncoded

# Bi directional recurrent neural network for NER



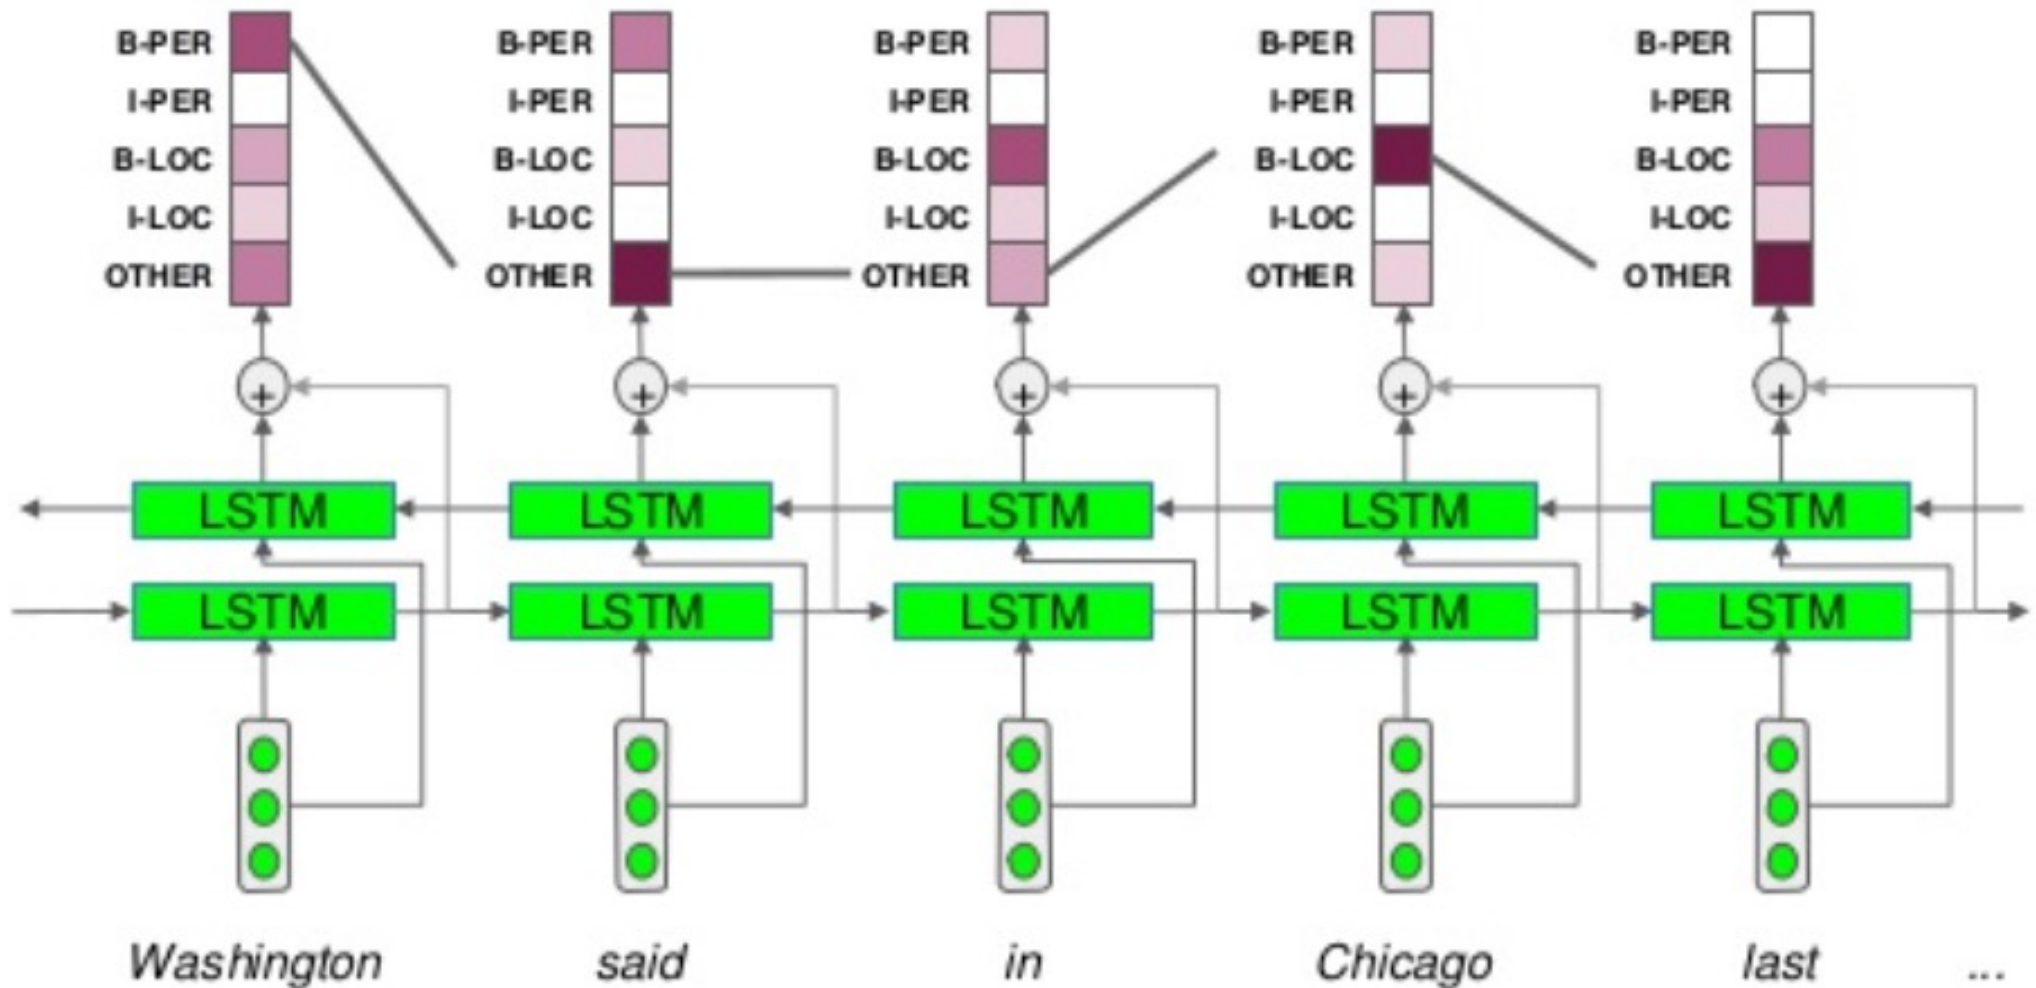
- Activation function for output: softmax
- Labels are OneHotEncoded

# Stacked Bi-RNN



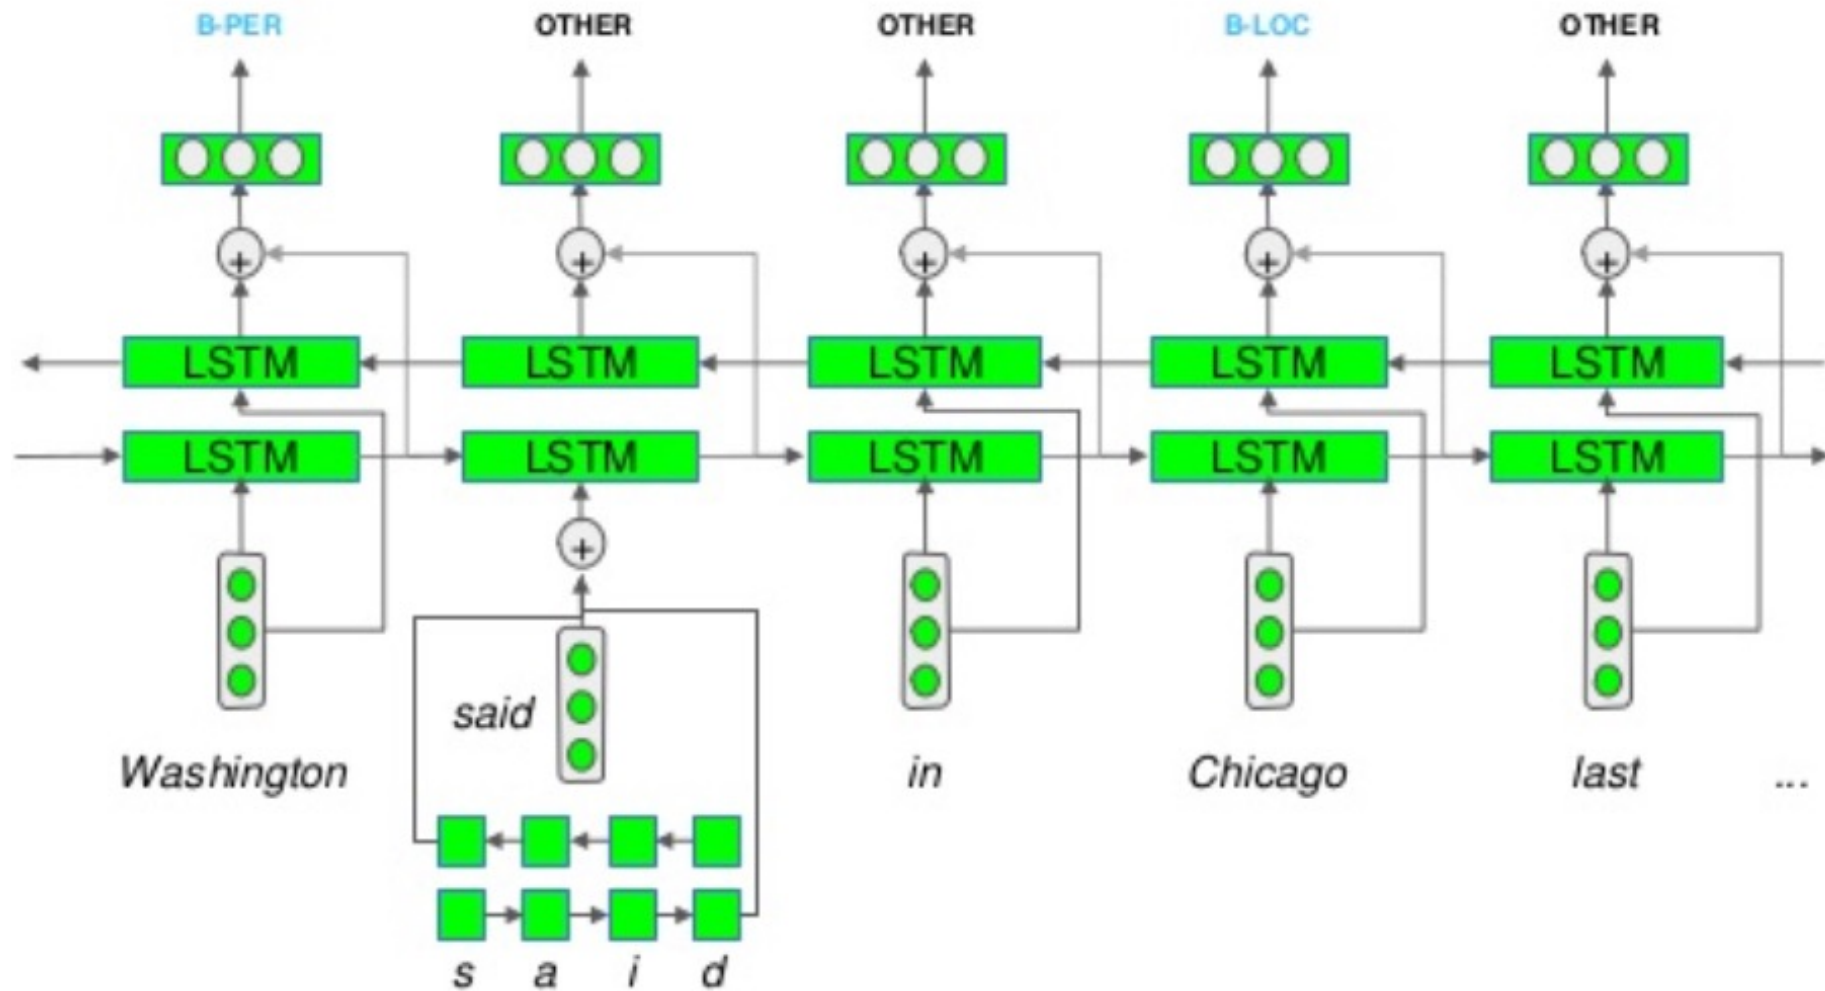
- Activation function for output: softmax
- Labels are OneHotEncoded

# Bi-RNN + CRF



- CRF output activation function

# Multi-level encoding char encoding + word encoding






# Other use of RNN

## → Sequence2sequence model

- ▶ Used for
  - ▶ Translation
  - ▶ Chatbot






# Some use of RNN

## → Input – Output Scenarios

Single - Single




One input

One output

→ Feed-forward network

Single -  
Multiple




One input

Many output

Image annotation

Multiple -  
Single




Many input

Many output

Text classification / Sentiment analysis

Multiple –  
Multiple



Many input / Many output  
 $\text{length}(\text{input}) \neq \text{length}(\text{output})$

Translation

Chat bot

Many input / Many output  
 $\text{length}(\text{input}) = \text{length}(\text{output})$

Named Entity Recognition

Part-Of-Speech tagging

# Other Useful Resources / References

---

- ▶ [http://cs231n.stanford.edu/slides/winter1516\\_lecture10.pdf](http://cs231n.stanford.edu/slides/winter1516_lecture10.pdf)
- ▶ <http://www.cs.toronto.edu/~rgrosse/csc321/lec10.pdf>
- ▶ R. Pascanu, T. Mikolov, and Y. Bengio, [On the difficulty of training recurrent neural networks](#), ICML 2013
- ▶ S. Hochreiter, and J. Schmidhuber, [Long short-term memory](#), Neural computation, 1997 9(8), pp.1735-1780
- ▶ F.A. Gers, and J. Schmidhuber, [Recurrent nets that time and count](#), IJCNN 2000
- ▶ K. Greff, R.K. Srivastava, J. Koutník, B.R. Steunebrink, and J. Schmidhuber, [LSTM: A search space odyssey](#), IEEE transactions on neural networks and learning systems, 2016
- ▶ K. Cho, B. Van Merriënboer, C. Gulcehre, D. Bahdanau, F. Bougares, H. Schwenk, and Y. Bengio, [Learning phrase representations using RNN encoder-decoder for statistical machine translation](#), ACL 2014
- ▶ R. Jozefowicz, W. Zaremba, and I. Sutskever, [An empirical exploration of recurrent network architectures](#), JMLR 2015