

Lab1: introduction to MPI

September 26, 2022

1 Grid'5000

For our lab sessions, we will use the Grid'5000 platform (<https://www.grid5000.fr>), and more precisely one of the clusters of the Lille site (see: <https://www.grid5000.fr/w/Lille:Hardware>). Each cluster is composed of multiple (homogeneous) compute nodes, with possibly several CPUs (= processor chip) per node, and multiple cores per CPU (currently one or more dozen).

A first tutorial on Grid'5000 can be found here: https://www.grid5000.fr/w/Getting_Started

Please read the usage policy: <https://www.grid5000.fr/w/Grid5000:UsagePolicy>

The current state of use of the Lille clusters is available here:

<https://intranet.grid5000.fr/oar/Lille/monika.cgi>

Warning: the Grid'5000 access which has been granted to you is only valid for the HPC teaching unit. **Any other non-academic usage is forbidden.** It is possible to use Grid'5000 compute resources for other academic work (project, internship ...), but you must warn me first (since I am your account manager).

1.1 Remote access

1.1.1 From a lab computer

In order to access the Lille site in Grid'5000, you must first obtain a ssh access to the outside of the lab room by requesting a session on: <https://vpn.centraledelille.fr>

Then, please write these lines in the `~/.ssh/config` file **on your local account** (create this file if it does not exist), replacing USERNAME by your Grid'5000 login:

```
Host_g5k
  User_USERNAME
  Hostname_access.grid5000.fr
  ForwardAgent_no

Host_*.g5k
  User_USERNAME
  ProxyCommand_ssh_g5k_W "$ (basename_%h_.g5k) :%p"
  ForwardAgent_no
```

You can now check that you can access Grid'5000 by running¹: (from a computer where your ssh key is stored) in a terminal (hereafter referred to as terminal 1):

```
$ ssh_USERNAME@lille.g5k
```

You should now be logged on a computer named `flille`.

1.1.2 From your laptop

On a Windows laptop, to have an ssh access with PuTTY or OpenSSH see here:

https://www.grid5000.fr/w/SSH#Windows_users

Then you should access Grid'5000 with something similar to (depending on your ssh client):

```
ssh -t_votre_login@access.grid5000.fr_ssh_lille
```

On a Linux laptop, see Section 1.1.1 (the VPN authorization may not be required).

¹The `$` character indicates the shell prompt.

1.2 Remote file editing

1.2.1 From a lab computer

For editing files on Grid’5000, we will mount the Grid’5000 file system through `ssh` with the following commands (to be run on your local computer, in another terminal – hereafter referred to as terminal 2):

```
$_chmod_755_$HOME
$_mkdir_G5K_lille
$_sshfs_-o_idmap=user_USERNAME@lille.g5k:/home/USERNAME_G5K_lille
```

You can now check that you can create directories and edit files on your local computer in the `G5K_lille` directory (in terminal 2), and that these are actually stored on the `flille` computer (`ls` and `cat` commands in terminal 1).

Important (possible file corruption otherwise!): at the end of the lab session, you have to “unmount” the file system mounted through `ssh` in `G5K_lille` with

```
$_fusermount_-u_G5K_lille/
```

The next time you want to mount/unmount the Grid’5000 file system through `ssh`, only the `sshfs/fusermount` commands are required.

1.2.2 From your laptop

On a Windows laptop, you should use an editor in text mode (`vi`, `vim`, `nano`...) in your `ssh` terminal to edit your source files.

On a Linux laptop, you can try `sshfs` as in Section 1.2.1, or `ssh` and an editor in text mode.

2 OpenMPI

On Grid’5000, we will use **OpenMPI** which implements the **MPI** standard. The programming mode is the **SPMD** mode: same executable for all processes, branchings depending on the process number (its **rank**) allowing them to perform different tasks.

To compile with MPI a C program named `foo.c` in an executable named `foo`, you have to use `mpicc` as:

```
$_mpicc_-o_foo_foo.c
```

With OpenMPI, one has to use the `mpirun` command to run this MPI program on a parallel architecture. However, within Grid’5000 the `mpirun` command has to be used with the `oarsub` and `oarsh` ones in order to obtain and access compute resources (CPU cores and memory) on one cluster. To do so, you must first configure OpenMPI **on your Grid’5000 account** (i.e. on `flille`) by adding these lines to the `~/ .openmpi/mca-params.conf` file (create the `~/ .openmpi` directory and this file if they do not exist):

```
plm_rsh_agent=oarsh
filem_rsh_agent=oarcp
btl_base_warn_component_unused=_0
```

We can now allocate resources (here 4 cores within 1 cluster) and run our MPI program on these resources thanks to the following command (which will create and run a “job”):

```
$_oarsub_-l_cluster=1/core=4'_mpirun_--mca_pml_^ucx_-machinefile_$OAR_NODEFILE_./foo'
```

`mpirun` will then use all allocated resources, and run one MPI process per core. For our labs, we will use one cluster at a time and only the requested number of cores (usually powers of 2 between 1 and 16) or the executable name should be changed. If one wants to provide arguments (e.g. `a1` and `a2`) on the command line of `foo`, these should be specified as:

```
$_oarsub_-l_cluster=1/core=4'_mpirun_--mca_pml_^ucx_-machinefile_$OAR_NODEFILE_./foo_a1_a2'
```

The cluster name can also be specified (in order to perform multiple performance tests on the same hardware) as:

```
$oarsub -l {"cluster='chetemi'"} /core=4 mpirun --mca_pml ^ucx -machinefile $OAR_NODEFILE ./foo'
```

If your job is long enough, you can see it executing by (re)loading the Monika web page (displaying the current state of use: <https://intranet.grid5000.fr/oar/Lille/monika.cgi>). Once your job is completed, you can check the output in a file named as `OAR.1234567.stdout` (where 1234567 is the job id). The error messages are printed in `OAR.1234567.stderr`. If your job does not start its execution, you can print informations on all submitted jobs with `oarstat` and delete your job with `oardel`.

For debugging your MPI programs, the easiest way (for our labs) is to rely on `printf` (taking care of the print order between different processes: see section 3).

3 Labs

You can start by downloading the latest MPI documentation (pdf format) at: <http://www.mpi-forum.org/>
For this first lab, we will start by using `Send` et `Recv` functions.

Exercise 1

What should be printed by the following MPI program?

Check this by writing, compiling and running this program on Grid'5000.

```
#include <stdio.h>
#include <mpi.h>

int main( int argc , char* argv[]){
    int rank , p , val , tag = 10;
    MPI_Status status;

    MPI_Init(&argc,&argv);
    MPI_Comm_size(MPI_COMM_WORLD, &p);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    if ( rank == 1){
        val = 18; MPI_Send(&val , 1 , MPI_INT , 0 , tag , MPI_COMM_WORLD);
    }
    else if ( rank == 0 ){
        MPI_Recv(&val , 1 , MPI_INT , 1 , tag , MPI_COMM_WORLD, &status);
        printf("I received the value %d from process with rank 1.\n",val);
    }

    MPI_Finalize();
}
```

Exercise 2

Here is a more advanced program:

```
#include <stdio.h>
#include <string.h>
#include <mpi.h>
#include <unistd.h>
#define SIZE_H_N 50

int main(int argc , char* argv[]){
    int my_rank; /* process rank */
```

```

int      p;                /* number of processes */
int      source;           /* sender rank          */
int      dest;             /* receiver rank        */
int      tag = 0;          /* message tag          */
char     msg[100];
MPI_Status status;
char     hostname[SIZE_H_N] ; gethostname(hostname, SIZE_H_N);

/* Initialization: */
MPI_Init(&argc, &argv);
MPI_Comm_rank(MPI_COMM_WORLD, &my_rank);
MPI_Comm_size(MPI_COMM_WORLD, &p);

/* Communications: */
if (my_rank != 0){
    /* Message creation */
    sprintf(msg, "Hello from process #%d from %s!", my_rank, hostname);
    dest = 0;

    MPI_Send(msg, strlen(msg)+1, MPI_CHAR, dest, tag, MPI_COMM_WORLD);
} else {
    for (source = 1; source < p; source++) {
        MPI_Recv(msg, 100, MPI_CHAR, source, tag, MPI_COMM_WORLD, &status);
        printf("On %s, the process #%d received the message: %s\n",
               hostname, my_rank, msg);
    }
}

/* Termination: */
MPI_Finalize();
}

```

question 1: same ordering of the displays
question 2: for a given process all the printf instructions appear on the 'correct ordering
question 3: The reception ordering varies from execution to the next
--> first message arrival means first message printed
question 4 : MPI_send >> MPI_Ssend
question 6: if the message size increase MPI Eager limit=ko

1. Run this MPI program multiple times with the same number of processes, and then by varying this number.
2. Add some `printf` at different places in the source code. What can you conclude?
3. Replace the `source` variable in the `MPI_Recv` call by the `MPI_ANY_SOURCE` wild card. Run some tests multiple times. What can you conclude?
4. Write an MPI program (in `ex_send.c`) where each process sends a string to its successor (i.e. to the process with rank `my_rank + 1` if `my_rank < p - 1`, and 0 otherwise), and receives a message from its predecessor. Once your program runs correctly, replace `MPI_Send` with `MPI_Ssend` in a new program (stored in `ex_ssend.c`). What is happening when executing this new program?
5. Copy `ex_ssend.c` in a new `ex_ssend_correct.c` file. In `ex_ssend_correct.c`, while still using `MPI_Ssend`, change the algorithm such that process 0 sends first its message to process 1, which will send its message to process 2 only after having received the message from process 0. In the same way, process 2 will send its message to process 3 only after having received the message from process 1, and so on...
6. In the original `ex_send.c` file, gradually increase the size of the data sent with `MPI_Send` until 100 KB. What is happening?

Exercise 3 – (optional)

Considering P processes, we want to implement a reduction algorithm which sums integer values from all processes (one integer per process). The root of the reduction algorithm will be the process with rank 0 (and the final result will thus be printed by this process only).

1. Design an efficient algorithm *without* any collective communication routine.
2. Write a parallel MPI program implementing such algorithm. The program will compute the global sum of random integer values generated by each process, such that the process with rank 0 will receive and print this global sum.
3. Rewrite your MPI program with a collective communication routine, and then modify it such that the global sum is available in all processes.